

Automated Trading with Cashew Capital

Report on an AI Studio Project

Advisors

Raisa Lobo, COO

Raisa@cashewcapitalfund.com

Bharath Venkataraman, CEO and CIO

Bharath@cashewcapitalfund.com

Boshen Parthasarathy, CTO

Boshen@cashewcapitalfund.com

Kevin McGee, CFO

Kevin@cashewcapitalfund.com

Contributors

Devina Gera

Fridah Ntikah

Humaira Sarwary

Kevin McGee

Krishna Kanodia

Lorraine Estrella De Leon

Lydia Chen

Nha-Van Nguyen

Osasikemwen Ogieva – oogieva25@amherst.edu

Source Code

<https://github.com/CashewCapital20/Automated-Trading/>

Table Of Contents

[Overview](#)

[Data](#)

[Training](#)

[Model Comparison](#)

[Real Time Trading](#)

[Back Testing](#)

[User Interfacing](#)

Overview

The goal of this project is to use AI and ML algorithms to make informed investment decisions by identifying trading opportunities using technical analysis.

We utilize data engineering, machine learning, and trading strategy evaluation to create a robust framework for predictive stock trading. Historical and real-time price data are sourced from Benzinga and persisted in MongoDB Atlas for efficient reuse.

Buy and sell signals are predetermined via percent changes between current and future prices and classification models are trained to predict them using technical indicators like EMA, RSI, MACD, volatility and Stochastic Oscillators.

Machine learning models we build include random forest, logistic regression, gradient boosting and ensemble stacking classifiers. We tune their hyperparameters with grids and evaluate them via classic metrics.

Real-time trading scripts simulate decision-making using delayed data from Benzinga and one of our pre-trained and persisted models. A user-friendly web interface collects user input in the form of a stock ticker and provides actionable buy, sell, or hold recommendations.

Data

Sources

Data was sourced primarily from [Benzinga](#), using its Python client. Some placeholder data was simulated.

Persistence

Whenever we fetch data afresh from Benzinga, to avoid hitting limits, we also upload it to a [MongoDB Atlas database](#), so we can continually reuse it.

This is done in the `fetch_load_mongo.py` file. This program takes in a stock symbol, fetches the last 383 days of historical price data for the asset as a list of price candles (open-high-low-close data), which is then converted to a dataframe and uploaded to MongoDB. It also uploads the last 26 days of data separately for recent analysis.

This file also provides functionality for calculating the technical indicators. These include the Exponential Moving Averages (the 12-period EMA, the 26-period EMA); the Relative Strength Index (RSI); the Stochastic Oscillator (the K and D lines); and the Moving Average Convergence Divergence (MACD), including its signal line and histogram. We then combine these indicators into the dataframe, our return value.

Throughout, there is careful exception handling and specified error messages for possible failures in code and the API.

The `test_fetch_load_mongo.py` file verifies that the above described data fetcher works as expected.

The `mongo_workflow_fix.py` file specifies a program which connects to our MongoDB Atlas database instance and retrieves documents we stored there earlier from Benzinga, based on a specific stock ticker.

Exploratory Analysis

The `dataset_exploratory_analysis.py` file fetches historical stock data from Benzinga's API for five stock tickers ("GOOG", "NVDA", "PG", "SPY", "PLTR") for the last 383 days.

Here we import and use the `calculate_indicators` function from our data fetcher and calculate the indicators on our collected data. Additionally, we calculate the rolling mean

and the 20-day rolling standard deviation, and visualize how all of these measures evolve with time, especially in comparison to the close price.

Plots are saved as images in the `/exploratory_analysis` folder.



Moving Averages and Close Price for GOOG

We also check for null values.

Training

The bulk of this work is in the `train_model.py` and `train_more_models.py` files.

Data, Labels and Features

We begin by connecting to our MongoDB collection, retrieving price data documents and converting them to a dataframe (in the `fetch_data` function).

The `prepare_labels` function generates a trading signal column based on future price movements of the asset. These signals will be used as the labels for training and evaluating our model.

Inputs into this function are `future_period` – the number of days ahead to look for the price change and `threshold` – the percentage change threshold to classify signals. Default values are `future_period=2` days, and `threshold=5%`.

We calculate a `future_price` column by shifting the price column upward by future period (2) rows. This allows comparison of the current close price to the future price, which lets us compute the `percentage_change` between the current price and the future price.

Finally, we use the percent change and the threshold to determine the signals. We start by initializing the `signal` column with all values set to 0.

- Set `signal` to 1 if the price change is greater than the positive threshold (indicating a "buy").
- Sets `signal` to -1 if the price decreases after a previous "buy" signal, indicating a "sell".
- Sets `signal` to -2 if the price change is less than the negative threshold (indicating a significant drop and a strong sell).
- Sets `signal` to 2 if the price increases after a previous "strong sell" signal, indicating a "re-entry buy".

The `prepare_features_and_labels` function segments our data into X (features) and y (labels) in preparation for model fitting. The features are the `macd`, the `macd_hist`, the `rsi`, the `slowk` and the `slowd` indicator columns. The newly created `signal` column is the label.

Random Forest Classifier Model

Our random forest classifier has 100 estimators, 20 maximum features and a maximum depth of 7. It was trained on 70% of our data with 30% reserved for testing, and saved as a `joblib` file in our repository.

Ensemble Model and Some Other Models

In `train_more_models.py`, we go on to build a logistic regression model, a support vector classifier, an AdaBoost classifier (with 100 estimators) and finally a stacking classifier.

The stacking classifier contains: a random forest classifier (with 100 estimators), a gradient boosting classifier (with 100 estimators), a logistic regression model, a support vector classifier and an AdaBoost with 100 estimators. It has a logistic regression model as the final estimator. This program is primarily about laying out a data pipeline for model training, evaluation and hyperparameter selection with Grid Search.

As such, the data is simulated. We generate 500 data points for our features (`MACD`, `RSI`, `Stochastic Oscillator`, and `close` price) from random distributions. We also include new indicators (`momentum`, `daily return`, and `volatility`). We also run some visualizations for these values to understand what our random data appears like. Then we define our `signal` column as a buy (1) or hold/sell (0) signal based on the next day's close price. We split data into training and testing sets (80% training, 20% testing).

After training, we evaluate each model using accuracy, a confusion matrix and a classification report (which provides detailed metrics like precision, recall, and F1-scores).

We use `GridSearchCV` for tuning hyperparameters (number of estimators and maximum depth) of both the random forest and gradient boosting models, logging and using the best parameters. Then we take a hybrid approach, predicting based on the averages of both these models. We evaluate this prediction, too, using accuracy, a confusion matrix and a classification report.

Model Comparison

The `model_comparison.py` script provides a comparative analysis of different machine learning models for stock trading predictions:

- **Regression Models:** Predict continuous outcomes and are evaluated with RMSE.
- **Classification Models:** Predict discrete outcomes (e.g., buy/sell signals) and are evaluated with accuracy.

Real Time Trading

Our implementation of a real-time trading system is an integration of several Python scripts.

The `delayed_stock_quote.py` fetches trading data from Benzinga which is 15 to 20 minutes behind the present, and this is what we use for simulating our real time trading.

The `real_time_trading.py` script executes real-time trading using this data and the pre-trained random forest classifier model we persisted as a .joblib file in our repository to make trading decisions.

These trade logs are uploaded to our MongoDB database by the `trade_logs.py` script. This program also calculates the remaining funds and inventory after each transaction.

The `paper_trade_demo.py` script is a stand-alone trading simulation which also uses our pre-trained random forest classifier model. It fetches random historical 5-minute candle data for a trading day, loads our model and simulates trading decisions based on model predictions, printing the results.

Back Testing

The script `back_test_predictions.py` is designed to evaluate the performance of a trading strategy by calculating various financial metrics, such as Sharpe Ratio, Max Drawdown, Profit Factor, and Win Rate. It also visualizes the cumulative returns and drawdowns over time.

User Interfacing

The `interface.py` and `app.py` scripts create a web-based, user-friendly interface for our project. Users should be able to enter a stock ticker and get a recommendation to buy, sell or hold.